

2b.Timing

November 2, 2024

1 Timing and optimizing code

Q: Are computers so fast that optimizing code for speed no longer matters?

1.1 The Context

- It depends
- Often optimizing isn't important
- Computers are faster
- It still can matter
- Try to learn & use best-practices

1.2 Video

[Modern Data Warehousing with BigQuery](#) - Watch from 5:00 to 7:20.

1.2.1 Take-aways

- You might care about speed when dealing with big data
- Some of the cutting edge problems to be solved *still* have to do with optimization

Considerations - Size of dataset - How often a task needs to be done - Time it takes to make process faster/time it takes to run it as is. - How much are you going to learn from the process of making your code faster?

- Is anybody else going to see your code?

1.3 How to speed up your code?

Below are a few easy ways.

1.3.1 1. List Comprehensions vs For Loop

```
[ ]: import numpy as np
```

```
[ ]: r = 1_000_000
```

```
[3]: # For Loop  
# %% will turn this into a cell magic  
# It will work on multi-line code  
%%timeit
```

```
squares = []

for n in range(r):
    squares.append(n**2)
```

507 ms \pm 113 ms per loop (mean \pm std. dev. of 7 runs, 1 loop each)

```
[4]: # List Comprehension
      # % indicates a magic command
      # %timeit will apply to one line
      result = %timeit -r3 -n10 -o squares = [n**2 for n in range(r)]
      result
```

380 ms \pm 63.5 ms per loop (mean \pm std. dev. of 3 runs, 10 loops each)

```
[4]: <TimeitResult : 380 ms  $\pm$  63.5 ms per loop (mean  $\pm$  std. dev. of 3 runs, 10 loops
each)>
```

```
[5]: result.all_runs
```

```
[5]: [4.687220814, 3.4598238140000035, 3.2474136390000012]
```

1.3.2 2. Numpy Arrays

```
[6]: %timeit squares = (np.arange(r))**2
```

1.8 ms \pm 172 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

```
[7]: %timeit squares = list((np.arange(r))**2)
```

63.6 ms \pm 11.1 ms per loop (mean \pm std. dev. of 7 runs, 10 loops each)

```
[8]: %timeit?
```

1.3.3 3. Multiple Assignments

```
[9]: %%timeit
      a = 5
      b = 10
      c = 20
      d = 25
```

34.7 ns \pm 0.611 ns per loop (mean \pm std. dev. of 7 runs, 10000000 loops each)

```
[10]: %timeit a,b,c,d = 5,10,20,25
```

34.3 ns \pm 0.798 ns per loop (mean \pm std. dev. of 7 runs, 10000000 loops each)

1.3.4 4. Use F Strings

```
[11]: name = 'Dwight Schrute'
```

```
[12]: %timeit name + ' loves beets'
```

97 ns ± 28.6 ns per loop (mean ± std. dev. of 7 runs, 10000000 loops each)

```
[13]: %timeit '%s loves beets' % name
```

168 ns ± 42.2 ns per loop (mean ± std. dev. of 7 runs, 10000000 loops each)

```
[14]: %timeit '{} loves beets'.format(name)
```

227 ns ± 8.8 ns per loop (mean ± std. dev. of 7 runs, 1000000 loops each)

```
[15]: %timeit f'{name} loves beets'
```

101 ns ± 27.4 ns per loop (mean ± std. dev. of 7 runs, 10000000 loops each)

1.3.5 5. Use Built-In Functions

Built-in functions such as `len()`, `abs()`, `min()`, and `max()` are written in C and are therefore very efficient.

1.3.6 6. Use Enumerate

```
[16]: our_list = [i**2 for i in range(5)]  
our_list
```

```
[16]: [0, 1, 4, 9, 16]
```

```
[17]: # Without enumerate  
# note using time only runs the script once  
%%time  
output = list(range(len(our_list)))  
for i in range(len(our_list)):  
    output[i] = (f'{i}^2 = {our_list[i]}')  
  
print("\n".join(output))
```

0² = 0

1² = 1

2² = 4

3² = 9

4² = 16

CPU times: user 128 µs, sys: 4 µs, total: 132 µs

Wall time: 391 µs

```
[18]: %%time  
for i, item in enumerate(our_list):
```

```
print(f'{i}^2 = {item}')
```

0² = 0

1² = 1

2² = 4

3² = 9

4² = 16

CPU times: user 126 µs, sys: 3 µs, total: 129 µs

Wall time: 134 µs

```
[19]: %%time
print("\n".join( f'{i}^2 = {item}' for i, item in enumerate(our_list) ) )
```

0² = 0

1² = 1

2² = 4

3² = 9

4² = 16

CPU times: user 1.85 ms, sys: 0 ns, total: 1.85 ms

Wall time: 1.76 ms

There are many other ways to speed up your code, you just have to decide if it is needed for whatever application you are working on.

1.4 Metric Prefixes & Scientific Notation

[Wikipedia Metric Prefixes & Symbols](#)

288.0 ms 10E-3 - decimal place goes over three -> 0.288 s

3.59 µs 10E-6 - decimal place goes over six places -> 0.00000359 s

[]: