

6b-Hierarchical.Clustering

November 13, 2024

1 Hierarchical Clustering.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
import scipy.cluster.hierarchy as shc
```

Load the data.

```
[2]: url = "https://ddc-datascience.s3.amazonaws.com/Wholesale_Data.csv"
data = pd.read_csv( url )
data.head()
```

```
[2]:
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844
3	1	3	13265	1196	4221	6404	507	1788
4	2	3	22615	5410	7198	3915	1777	5185

```
[3]: data.shape
```

```
[3]: (440, 8)
```

```
[4]: data.describe().transpose()
```

```
[4]:
```

	count	mean	std	min	25%	50%	\
Channel	440.0	1.322727	0.468052	1.0	1.00	1.0	
Region	440.0	2.543182	0.774272	1.0	2.00	3.0	
Fresh	440.0	12000.297727	12647.328865	3.0	3127.75	8504.0	
Milk	440.0	5796.265909	7380.377175	55.0	1533.00	3627.0	
Grocery	440.0	7951.277273	9503.162829	3.0	2153.00	4755.5	
Frozen	440.0	3071.931818	4854.673333	25.0	742.25	1526.0	
Detergents_Paper	440.0	2881.493182	4767.854448	3.0	256.75	816.5	
Delicassen	440.0	1524.870455	2820.105937	3.0	408.25	965.5	

	75%	max
Channel	2.00	2.0
Region	3.00	3.0
Fresh	16933.75	112151.0
Milk	7190.25	73498.0
Grocery	10655.75	92780.0
Frozen	3554.25	60869.0
Detergents_Paper	3922.00	40827.0
Delicassen	1820.25	47943.0

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 440 entries, 0 to 439
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Channel              440 non-null    int64
1   Region               440 non-null    int64
2   Fresh                440 non-null    int64
3   Milk                 440 non-null    int64
4   Grocery              440 non-null    int64
5   Frozen               440 non-null    int64
6   Detergents_Paper     440 non-null    int64
7   Delicassen           440 non-null    int64
dtypes: int64(8)
memory usage: 27.6 KB
```

```
[6]: data[["Channel", "Region"]].nunique()
```

```
[6]: Channel    2
      Region    3
      dtype: int64
```

```
[8]: data[["Channel", "Region"]].value_counts( )
```

```
[8]: Channel  Region
1         3      211
2         3      105
1         1       59
         2       28
2         2       19
         1       18
Name: count, dtype: int64
```

Scale the data.

```
[9]: # Scale data
scaler = MinMaxScaler()
scaler.fit(data)
data_scaled = scaler.transform(data)
# Convert back to data frame
data_scaled = pd.DataFrame(data_scaled, columns = data.columns)
data_scaled.head()
```

```
[9]:
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	\
0	1.0	1.0	0.112940	0.130727	0.081464	0.003106	0.065427	
1	1.0	1.0	0.062899	0.132824	0.103097	0.028548	0.080590	
2	1.0	1.0	0.056622	0.119181	0.082790	0.039116	0.086052	
3	0.0	1.0	0.118254	0.015536	0.045464	0.104842	0.012346	
4	1.0	1.0	0.201626	0.072914	0.077552	0.063934	0.043455	

	Delicassen
0	0.027847
1	0.036984
2	0.163559
3	0.037234
4	0.108093

```
[10]: data_scaled.describe().transpose()
```

```
[10]:
```


	count	mean	std	min	25%	50%	\
Channel	440.0	0.322727	0.468052	0.0	0.000000	0.000000	
Region	440.0	0.771591	0.387136	0.0	0.500000	1.000000	
Fresh	440.0	0.106977	0.112774	0.0	0.027863	0.075802	
Milk	440.0	0.078173	0.100491	0.0	0.020124	0.048636	
Grocery	440.0	0.085671	0.102430	0.0	0.023174	0.051225	
Frozen	440.0	0.050078	0.079789	0.0	0.011788	0.024670	
Detergents_Paper	440.0	0.070510	0.116790	0.0	0.006216	0.019927	
Delicassen	440.0	0.031745	0.058826	0.0	0.008453	0.020077	

	75%	max
Channel	1.000000	1.0
Region	1.000000	1.0
Fresh	0.150968	1.0
Milk	0.097154	1.0
Grocery	0.114821	1.0
Frozen	0.058005	1.0
Detergents_Paper	0.095997	1.0
Delicassen	0.037907	1.0

Create dendrogram of the clustering.


```
[11]: plt.figure(figsize=(10, 7))
plt.title("Dendrograms")
```

```
dend = shc.dendrogram(shc.linkage(data_scaled, method='ward'))
```



Choose an appropriate threshold for the clusters.

```
[12]: plt.figure(figsize=(10, 7))  
plt.title("Dendrograms")  
dend = shc.dendrogram(shc.linkage(data_scaled, method='ward'))  
plt.axhline(y=6, color='r', linestyle='--') ;
```



Classify all points based on the number of clusters at the threshold you chose.

```
[13]: from sklearn.cluster import AgglomerativeClustering
      cluster = AgglomerativeClustering(n_clusters=3, metric='euclidean',
      linkage='ward')
      clusterNums = pd.Series(cluster.fit_predict(data_scaled))
```

```
[14]: clusterNums
```

```
[14]: 0      0
      1      0
      2      0
      3      2
      4      0
      ..
      435    2
      436    2
      437    0
      438    2
      439    2
      Length: 440, dtype: int64
```

```
[15]: clusterNums.value_counts()
```

```
[15]: 2    212
      0    142
      1     86
      Name: count, dtype: int64
```

Visualize the clusters.

```
[16]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 440 entries, 0 to 439
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Channel               440 non-null    int64
1   Region                440 non-null    int64
2   Fresh                 440 non-null    int64
3   Milk                  440 non-null    int64
4   Grocery               440 non-null    int64
5   Frozen                440 non-null    int64
6   Detergents_Paper      440 non-null    int64
7   Delicassen            440 non-null    int64
dtypes: int64(8)
memory usage: 27.6 KB
```

```
[17]: cluster.labels_
```

```
[17]: array([0, 0, 0, 2, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 2, 0, 2, 0, 2, 0, 2,
        2, 0, 0, 0, 2, 2, 0, 2, 2, 2, 2, 2, 2, 0, 2, 0, 0, 2, 2, 2, 0, 0,
        0, 0, 0, 0, 0, 0, 2, 2, 0, 0, 2, 2, 0, 0, 2, 2, 0, 0, 0, 0, 2, 0,
        2, 0, 2, 2, 2, 2, 2, 0, 0, 2, 2, 0, 2, 2, 2, 0, 0, 2, 0, 0, 0, 2,
        2, 2, 2, 2, 0, 2, 0, 2, 0, 2, 2, 2, 0, 0, 0, 2, 2, 2, 0, 0, 0, 0,
        2, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 2, 2, 2, 0, 2, 2, 2, 2,
        2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 2, 2, 2, 2, 2, 2, 2,
        2, 0, 0, 2, 0, 0, 0, 2, 2, 0, 0, 0, 0, 2, 2, 2, 0, 0, 2, 0, 2, 0,
        2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 0, 2, 2, 2, 0, 2, 2, 1, 0,
        1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1,
        1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        0, 1, 0, 1, 0, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 0, 2, 0, 2, 2, 2, 2,
        2, 2, 2, 2, 2, 2, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
        1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 2, 1, 1, 1, 1,
        1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 2, 0, 2, 2, 0, 0, 2, 0, 2, 0,
        2, 0, 2, 2, 2, 0, 2, 2, 2, 2, 2, 2, 2, 0, 2, 2, 2, 2, 0, 2, 2, 0,
        2, 2, 0, 2, 2, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
        0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 0, 2, 2, 2, 2, 2, 2, 0, 0, 2,
```


```
0, 2, 2, 0, 2, 0, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 0, 2, 2])
```

```
[18]: data['clusters'] = cluster.labels_
```

```
[19]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 440 entries, 0 to 439
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Channel                440 non-null    int64
1   Region                 440 non-null    int64
2   Fresh                  440 non-null    int64
3   Milk                   440 non-null    int64
4   Grocery                440 non-null    int64
5   Frozen                 440 non-null    int64
6   Detergents_Paper       440 non-null    int64
7   Delicassen             440 non-null    int64
8   clusters               440 non-null    int64
dtypes: int64(9)
memory usage: 31.1 KB
```

```
[20]: sns.pairplot(data, hue = 'clusters') ;
```



[]: