

03.07-Merge-and-Join

November 10, 2024

1 Combining Datasets: merge and join

One important feature offered by Pandas is its high-performance, in-memory join and merge operations, which you may be familiar with if you have ever worked with databases. The main interface for this is the `pd.merge` function, and we'll see a few examples of how this can work in practice.

For convenience, we will again define the `display` function from the previous chapter after the usual imports:

```
[1]: import pandas as pd
import numpy as np

class display(object):
    """Display HTML representation of multiple objects"""
    template = """<div style="float: left; padding: 10px;">
<p style='font-family:"Courier New", Courier, monospace'>{0}</p>{1}
</div>"""
    def __init__(self, *args):
        self.args = args

    def _repr_html_(self):
        return '\n'.join(self.template.format(a, eval(a)._repr_html_())
                          for a in self.args)

    def __repr__(self):
        return '\n\n'.join(a + '\n' + repr(eval(a))
                           for a in self.args)
```

1.1 Relational Algebra

The behavior implemented in `pd.merge` is a subset of what is known as *relational algebra*, which is a formal set of rules for manipulating relational data that forms the conceptual foundation of operations available in most databases. The strength of the relational algebra approach is that it proposes several fundamental operations, which become the building blocks of more complicated operations on any dataset. With this lexicon of fundamental operations implemented efficiently in a database or other program, a wide range of fairly complicated composite operations can be performed.

Pandas implements several of these fundamental building blocks in the `pd.merge` function and the

related `join` method of `Series` and `DataFrame` objects. As you will see, these let you efficiently link data from different sources.

1.2 Categories of Joins

The `pd.merge` function implements a number of types of joins: *one-to-one*, *many-to-one*, and *many-to-many*. All three types of joins are accessed via an identical call to the `pd.merge` interface; the type of join performed depends on the form of the input data. We'll start with some simple examples of the three types of merges, and discuss detailed options a bit later.

1.2.1 One-to-One Joins

Perhaps the simplest type of merge is the one-to-one join, which is in many ways similar to the column-wise concatenation you saw in [Combining Datasets: Concat & Append](#). As a concrete example, consider the following two `DataFrame` objects, which contain information on several employees in a company:

```
[2]: df1 = pd.DataFrame({'employee': ['Bob', 'Jake', 'Lisa', 'Sue'],
                        'group': ['Accounting', 'Engineering',
                                'Engineering', 'HR']})
df2 = pd.DataFrame({'employee': ['Lisa', 'Bob', 'Jake', 'Sue'],
                    'hire_date': [2004, 2008, 2012, 2014]})
display('df1', 'df2')
```

```
[2]: df1
   employee  group
0      Bob  Accounting
1      Jake  Engineering
2      Lisa  Engineering
3       Sue         HR

df2
   employee  hire_date
0      Lisa      2004
1       Bob      2008
2      Jake      2012
3       Sue      2014
```

To combine this information into a single `DataFrame`, we can use the `pd.merge` function:

```
[3]: df3 = pd.merge(df1, df2)
df3
```

```
[3]:   employee  group  hire_date
0      Bob  Accounting      2008
1      Jake  Engineering      2012
2      Lisa  Engineering      2004
3       Sue         HR      2014
```

The `pd.merge` function recognizes that each `DataFrame` has an `employee` column, and automatically joins using this column as a key. The result of the merge is a new `DataFrame` that combines the information from the two inputs. Notice that the order of entries in each column is not necessarily maintained: in this case, the order of the `employee` column differs between `df1` and `df2`, and the `pd.merge` function correctly accounts for this. Additionally, keep in mind that the merge in general discards the index, except in the special case of merges by index (see the `left_index` and `right_index` keywords, discussed momentarily).

1.2.2 Many-to-One Joins

Many-to-one joins are joins in which one of the two key columns contains duplicate entries. For the many-to-one case, the resulting `DataFrame` will preserve those duplicate entries as appropriate. Consider the following example of a many-to-one join:

```
[4]: df4 = pd.DataFrame({'group': ['Accounting', 'Engineering', 'HR'],
                        'supervisor': ['Carly', 'Guido', 'Steve']})
display('df3', 'df4', 'pd.merge(df3, df4)')
```

```
[4]: df3
   employee  group  hire_date
0      Bob  Accounting      2008
1     Jake  Engineering      2012
2     Lisa  Engineering      2004
3      Sue         HR       2014

df4
   group supervisor
0  Accounting    Carly
1  Engineering    Guido
2         HR      Steve

pd.merge(df3, df4)
   employee  group  hire_date  supervisor
0      Bob  Accounting      2008      Carly
1     Jake  Engineering      2012      Guido
2     Lisa  Engineering      2004      Guido
3      Sue         HR       2014      Steve
```

The resulting `DataFrame` has an additional column with the “supervisor” information, where the information is repeated in one or more locations as required by the inputs.

1.2.3 Many-to-Many Joins

Many-to-many joins may be a bit confusing conceptually, but are nevertheless well defined. If the key column in both the left and right arrays contains duplicates, then the result is a many-to-many merge. This will be perhaps most clear with a concrete example. Consider the following, where we have a `DataFrame` showing one or more skills associated with a particular group. By performing a many-to-many join, we can recover the skills associated with any individual person:

```
[5]: df5 = pd.DataFrame({'group': ['Accounting', 'Accounting',
                                   'Engineering', 'Engineering', 'HR', 'HR'],
                        'skills': ['math', 'spreadsheets', 'software', 'math',
                                   'spreadsheets', 'organization']})
display('df1', 'df5', "pd.merge(df1, df5)")
```

```
[5]: df1
   employee  group
0      Bob  Accounting
1      Jake  Engineering
2      Lisa  Engineering
3       Sue           HR

df5
   group  skills
0  Accounting    math
1  Accounting spreadsheets
2  Engineering    software
3  Engineering    math
4           HR spreadsheets
5           HR  organization

pd.merge(df1, df5)
   employee  group  skills
0      Bob  Accounting    math
1      Bob  Accounting spreadsheets
2      Jake  Engineering    software
3      Jake  Engineering    math
4      Lisa  Engineering    software
5      Lisa  Engineering    math
6       Sue           HR spreadsheets
7       Sue           HR  organization
```

These three types of joins can be used with other Pandas tools to implement a wide array of functionality. But in practice, datasets are rarely as clean as the one we're working with here. In the following section we'll consider some of the options provided by `pd.merge` that enable you to tune how the join operations work.

1.3 Specification of the Merge Key

We've already seen the default behavior of `pd.merge`: it looks for one or more matching column names between the two inputs, and uses this as the key. However, often the column names will not match so nicely, and `pd.merge` provides a variety of options for handling this.

```
[ ]:
```

1.3.1 The on Keyword

Most simply, you can explicitly specify the name of the key column using the `on` keyword, which takes a column name or a list of column names:

```
[6]: display('df1', 'df2', "pd.merge(df1, df2, on='employee')")
```

```
[6]: df1
      employee      group
0      Bob  Accounting
1      Jake  Engineering
2      Lisa  Engineering
3      Sue      HR

df2
      employee  hire_date
0      Lisa      2004
1      Bob      2008
2      Jake      2012
3      Sue      2014

pd.merge(df1, df2, on='employee')
      employee      group  hire_date
0      Bob  Accounting      2008
1      Jake  Engineering      2012
2      Lisa  Engineering      2004
3      Sue      HR      2014
```

This option works only if both the left and right `DataFrames` have the specified column name.

1.3.2 The left_on and right_on Keywords

At times you may wish to merge two datasets with different column names; for example, we may have a dataset in which the employee name is labeled as “name” rather than “employee”. In this case, we can use the `left_on` and `right_on` keywords to specify the two column names:

```
[7]: df3 = pd.DataFrame({'name': ['Bob', 'Jake', 'Lisa', 'Sue'],
                        'salary': [70000, 80000, 120000, 90000]})
display('df1', 'df3', 'pd.merge(df1, df3, left_on="employee", right_on="name")')
```

```
[7]: df1
      employee      group
0      Bob  Accounting
1      Jake  Engineering
2      Lisa  Engineering
3      Sue      HR

df3
      name  salary
```

```

0   Bob   70000
1   Jake  80000
2   Lisa 120000
3   Sue   90000

```

```

pd.merge(df1, df3, left_on="employee", right_on="name")

```

	employee	group	name	salary
0	Bob	Accounting	Bob	70000
1	Jake	Engineering	Jake	80000
2	Lisa	Engineering	Lisa	120000
3	Sue	HR	Sue	90000

The result has a redundant column that we can drop if desired—for example, by using the `DataFrame.drop()` method:

```
[8]: pd.merge(df1, df3, left_on="employee", right_on="name").drop('name', axis=1)
```

```
[8]:
```

	employee	group	salary
0	Bob	Accounting	70000
1	Jake	Engineering	80000
2	Lisa	Engineering	120000
3	Sue	HR	90000

1.3.3 The `left_index` and `right_index` Keywords

Sometimes, rather than merging on a column, you would instead like to merge on an index. For example, your data might look like this:

```
[9]: df1a = df1.set_index('employee')
df2a = df2.set_index('employee')
display(df1a, df2a)
```

```
[9]: df1a
```

	group
employee	
Bob	Accounting
Jake	Engineering
Lisa	Engineering
Sue	HR


```
df2a
```

	hire_date
employee	
Lisa	2004
Bob	2008
Jake	2012
Sue	2014

You can use the index as the key for merging by specifying the `left_index` and/or `right_index` flags in `pd.merge()`:

```
[10]: display('df1a', 'df2a',  
            "pd.merge(df1a, df2a, left_index=True, right_index=True)")
```

```
[10]: df1a
```

	group
employee	
Bob	Accounting
Jake	Engineering
Lisa	Engineering
Sue	HR


```
df2a
```

	hire_date
employee	
Lisa	2004
Bob	2008
Jake	2012
Sue	2014


```
pd.merge(df1a, df2a, left_index=True, right_index=True)
```

	group	hire_date
employee		
Bob	Accounting	2008
Jake	Engineering	2012
Lisa	Engineering	2004
Sue	HR	2014

For convenience, Pandas includes the `DataFrame.join()` method, which performs an index-based merge without extra keywords:

```
[11]: df1a.join(df2a)
```

```
[11]:
```

	group	hire_date
employee		
Bob	Accounting	2008
Jake	Engineering	2012
Lisa	Engineering	2004
Sue	HR	2014

If you'd like to mix indices and columns, you can combine `left_index` with `right_on` or `left_on` with `right_index` to get the desired behavior:

```
[12]: display('df1a', 'df3', "pd.merge(df1a, df3, left_index=True, right_on='name')")
```

```
[12]: df1a
```

	group
--	-------

```

employee
Bob      Accounting
Jake      Engineering
Lisa      Engineering
Sue              HR

```

```

df3
   name  salary
0   Bob   70000
1  Jake   80000
2  Lisa  120000
3   Sue   90000

```

```

pd.merge(df1a, df3, left_index=True, right_on='name')
   group name  salary
0  Accounting  Bob   70000
1  Engineering Jake   80000
2  Engineering Lisa  120000
3         HR   Sue   90000

```

All of these options also work with multiple indices and/or multiple columns; the interface for this behavior is very intuitive. For more information on this, see the [“Merge, Join, and Concatenate” section](#) of the Pandas documentation.

1.4 Specifying Set Arithmetic for Joins

In all the preceding examples we have glossed over one important consideration in performing a join: the type of set arithmetic used in the join. This comes up when a value appears in one key column but not the other. Consider this example:

```

[13]: df6 = pd.DataFrame({'name': ['Peter', 'Paul', 'Mary'],
                        'food': ['fish', 'beans', 'bread']},
                        columns=['name', 'food'])
df7 = pd.DataFrame({'name': ['Mary', 'Joseph'],
                    'drink': ['wine', 'beer']},
                    columns=['name', 'drink'])
display('df6', 'df7', 'pd.merge(df6, df7)')

```

```

[13]: df6
   name  food
0 Peter  fish
1  Paul  beans
2  Mary  bread

df7
   name drink
0  Mary  wine
1 Joseph  beer

```



```
pd.merge(df6, df7)
   name  food drink
0  Mary  bread  wine
```

Here we have merged two datasets that have only a single “name” entry in common: Mary. By default, the result contains the *intersection* of the two sets of inputs; this is what is known as an *inner join*. We can specify this explicitly using the `how` keyword, which defaults to “inner”:

```
[14]: pd.merge(df6, df7, how='inner')
```

```
[14]:   name  food drink
      0  Mary  bread  wine
```

Other options for the `how` keyword are 'outer', 'left', and 'right'. An *outer join* returns a join over the union of the input columns, and fills in all missing values with NAs:

```
[15]: display('df6', 'df7', "pd.merge(df6, df7, how='outer')")
```

```
[15]: df6
      name  food
0  Peter  fish
1   Paul  beans
2   Mary  bread

df7
      name drink
0   Mary  wine
1  Joseph  beer

pd.merge(df6, df7, how='outer')
   name  food drink
0  Peter  fish  NaN
1   Paul  beans  NaN
2   Mary  bread  wine
3  Joseph   NaN  beer
```

The *left join* and *right join* return joins over the left entries and right entries, respectively. For example:

```
[16]: display('df6', 'df7', "pd.merge(df6, df7, how='left')")
```

```
[16]: df6
      name  food
0  Peter  fish
1   Paul  beans
2   Mary  bread

df7
```

```

      name drink
0   Mary  wine
1  Joseph  beer

```

```

pd.merge(df6, df7, how='left')
      name  food drink
0  Peter   fish  NaN
1   Paul  beans  NaN
2   Mary  bread  wine

```

The output rows now correspond to the entries in the left input. Using `how='right'` works in a similar manner.

All of these options can be applied straightforwardly to any of the preceding join types.

1.5 Overlapping Column Names: The suffixes Keyword

Last, you may end up in a case where your two input `DataFrames` have conflicting column names. Consider this example:

```

[17]: df8 = pd.DataFrame({'name': ['Bob', 'Jake', 'Lisa', 'Sue'],
                        'rank': [1, 2, 3, 4]})
df9 = pd.DataFrame({'name': ['Bob', 'Jake', 'Lisa', 'Sue'],
                        'rank': [3, 1, 4, 2]})
display('df8', 'df9', 'pd.merge(df8, df9, on="name")')

```

```

[17]: df8
      name  rank
0    Bob     1
1    Jake     2
2    Lisa     3
3    Sue     4

df9
      name  rank
0    Bob     3
1    Jake     1
2    Lisa     4
3    Sue     2

pd.merge(df8, df9, on="name")
      name  rank_x  rank_y
0    Bob         1         3
1    Jake         2         1
2    Lisa         3         4
3    Sue         4         2

```

Because the output would have two conflicting column names, the `merge` function automatically appends the suffixes `_x` and `_y` to make the output columns unique. If these defaults are inappro-

priate, it is possible to specify a custom suffix using the `suffixes` keyword:

```
[18]: pd.merge(df8, df9, on="name", suffixes=["_L", "_R"])
```

```
[18]:
```

	name	rank_L	rank_R
0	Bob	1	3
1	Jake	2	1
2	Lisa	3	4
3	Sue	4	2

These suffixes work in any of the possible join patterns, and also work if there are multiple overlapping columns.

For more information on these patterns, see [Aggregation and Grouping](#), where we dive a bit deeper into relational algebra. Also see the “[Merge, Join, Concatenate and Compare](#)” section of the Pandas documentation for further discussion of these topics.

1.6 Example: US States Data

Merge and join operations come up most often when combining data from different sources. Here we will consider an example of some data about US states and their populations. The data files can be found at <http://github.com/jakevdp/data-USstates>:

```
[19]: # Following are commands to download the data
# repo = "https://raw.githubusercontent.com/jakevdp/data-USstates/master"
# !cd data && curl -O {repo}/state-population.csv
# !cd data && curl -O {repo}/state-areas.csv
# !cd data && curl -O {repo}/state-abbrevs.csv
```

Let’s take a look at the three datasets, using the Pandas `read_csv` function:

```
[20]: pop = pd.read_csv('data/state-population.csv')
areas = pd.read_csv('data/state-areas.csv')
abbrevs = pd.read_csv('data/state-abbrevs.csv')

display('pop.head()', 'areas.head()', 'abbrevs.head()')
```

```
[20]: pop.head()
```

	state/region	ages	year	population
0	AL	under18	2012	1117489.0
1	AL	total	2012	4817528.0
2	AL	under18	2010	1130966.0
3	AL	total	2010	4785570.0
4	AL	under18	2011	1125763.0

```
areas.head()
```

	state	area (sq. mi)
0	Alabama	52423
1	Alaska	656425
2	Arizona	114006

```
3    Arkansas      53182
4    California   163707
```

```
abbrevs.head()
      state abbreviation
0    Alabama          AL
1    Alaska           AK
2    Arizona          AZ
3    Arkansas         AR
4    California       CA
```

Given this information, say we want to compute a relatively straightforward result: rank US states and territories by their 2010 population density. We clearly have the data here to find this result, but we'll have to combine the datasets to do so.

We'll start with a many-to-one merge that will give us the full state names within the population DataFrame. We want to merge based on the `state/region` column of `pop` and the `abbreviation` column of `abbrevs`. We'll use `how='outer'` to make sure no data is thrown away due to mismatched labels:

```
[21]: merged = pd.merge(pop, abbrevs, how='outer',
                        left_on='state/region', right_on='abbreviation')
merged = merged.drop('abbreviation', axis=1) # drop duplicate info
merged.head()
```

```
[21]:   state/region  ages  year  population  state
0         AL  under18  2012   1117489.0  Alabama
1         AL    total  2012   4817528.0  Alabama
2         AL  under18  2010   1130966.0  Alabama
3         AL    total  2010   4785570.0  Alabama
4         AL  under18  2011   1125763.0  Alabama
```

Let's double-check whether there were any mismatches here, which we can do by looking for rows with nulls:

```
[22]: merged.isnull().any()
```

```
[22]: state/region    False
ages              False
year              False
population         True
state             True
dtype: bool
```

Some of the population values are null; let's figure out which these are!

```
[23]: merged[merged['population'].isnull()].head()
```

```
[23]:
```

	state/region	ages	year	population	state
2448	PR	under18	1990	NaN	NaN
2449	PR	total	1990	NaN	NaN
2450	PR	total	1991	NaN	NaN
2451	PR	under18	1991	NaN	NaN
2452	PR	total	1993	NaN	NaN

It appears that all the null population values are from Puerto Rico prior to the year 2000; this is likely due to this data not being available in the original source.

More importantly, we see that some of the new `state` entries are also null, which means that there was no corresponding entry in the `abbrevs` key! Let's figure out which regions lack this match:

```
[24]: merged.loc[merged['state'].isnull(), 'state/region'].unique()
```

```
[24]: array(['PR', 'USA'], dtype=object)
```

We can quickly infer the issue: our population data includes entries for Puerto Rico (PR) and the United States as a whole (USA), while these entries do not appear in the state abbreviation key. We can fix these quickly by filling in appropriate entries:

```
[25]: merged.loc[merged['state/region'] == 'PR', 'state'] = 'Puerto Rico'
merged.loc[merged['state/region'] == 'USA', 'state'] = 'United States'
merged.isnull().any()
```

```
[25]: state/region    False
ages                False
year                False
population           True
state                False
dtype: bool
```

No more nulls in the `state` column: we're all set!

Now we can merge the result with the area data using a similar procedure. Examining our results, we will want to join on the `state` column in both:

```
[26]: final = pd.merge(merged, areas, on='state', how='left')
final.head()
```

```
[26]:
```

	state/region	ages	year	population	state	area (sq. mi)
0	AL	under18	2012	1117489.0	Alabama	52423.0
1	AL	total	2012	4817528.0	Alabama	52423.0
2	AL	under18	2010	1130966.0	Alabama	52423.0
3	AL	total	2010	4785570.0	Alabama	52423.0
4	AL	under18	2011	1125763.0	Alabama	52423.0

Again, let's check for nulls to see if there were any mismatches:

```
[27]: final.isnull().any()
```

```
[27]: state/region    False
      ages           False
      year           False
      population     True
      state          False
      area (sq. mi)  True
      dtype: bool
```

There are nulls in the `area` column; we can take a look to see which regions were ignored here:

```
[28]: final['state'][final['area (sq. mi)'].isnull()].unique()
```

```
[28]: array(['United States'], dtype=object)
```

We see that our `areas` `DataFrame` does not contain the area of the United States as a whole. We could insert the appropriate value (using the sum of all state areas, for instance), but in this case we'll just drop the null values because the population density of the entire United States is not relevant to our current discussion:

```
[29]: final.dropna(inplace=True)
      final.head()
```

```
[29]:   state/region  ages  year  population  state  area (sq. mi)
0         AL  under18  2012   1117489.0  Alabama    52423.0
1         AL    total  2012   4817528.0  Alabama    52423.0
2         AL  under18  2010   1130966.0  Alabama    52423.0
3         AL    total  2010   4785570.0  Alabama    52423.0
4         AL  under18  2011   1125763.0  Alabama    52423.0
```

Now we have all the data we need. To answer the question of interest, let's first select the portion of the data corresponding with the year 2010, and the total population. We'll use the `query` function to do this quickly (this requires the `NumExpr` package to be installed; see [High-Performance Pandas: `eval\(\)` and `query\(\)`](#)):

```
[30]: data2010 = final.query("year == 2010 & ages == 'total'")
      data2010.head()
```

```
[30]:   state/region  ages  year  population  state  area (sq. mi)
3         AL    total  2010   4785570.0  Alabama    52423.0
91        AK    total  2010    713868.0   Alaska   656425.0
101       AZ    total  2010   6408790.0   Arizona   114006.0
189       AR    total  2010   2922280.0  Arkansas    53182.0
197       CA    total  2010  37333601.0  California  163707.0
```

Now let's compute the population density and display it in order. We'll start by re-indexing our data on the state, and then compute the result:

```
[31]: data2010.set_index('state', inplace=True)
      density = data2010['population'] / data2010['area (sq. mi)']
```

```
[32]: density.sort_values(ascending=False, inplace=True)
density.head()
```

```
[32]: state
District of Columbia    8898.897059
Puerto Rico            1058.665149
New Jersey              1009.253268
Rhode Island            681.339159
Connecticut             645.600649
dtype: float64
```

The result is a ranking of US states, plus Washington, DC, and Puerto Rico, in order of their 2010 population density, in residents per square mile. We can see that by far the densest region in this dataset is Washington, DC (i.e., the District of Columbia); among states, the densest is New Jersey.

We can also check the end of the list:

```
[33]: density.tail()
```

```
[33]: state
South Dakota    10.583512
North Dakota    9.537565
Montana         6.736171
Wyoming         5.768079
Alaska          1.087509
dtype: float64
```

We see that the least dense state, by far, is Alaska, averaging slightly over one resident per square mile.

This type of data merging is a common task when trying to answer questions using real-world data sources. I hope that this example has given you an idea of some of the ways you can combine the tools we've covered in order to gain insight from your data!